


# Two-Lane Traffic Simulation


Tommy Poek

47192085

# Intro - Visualization



# Intro - Visualization



# Intro - Scaling

## ► Scale up #lanes



## ► Scale up #cars



Fair benchmarking ←

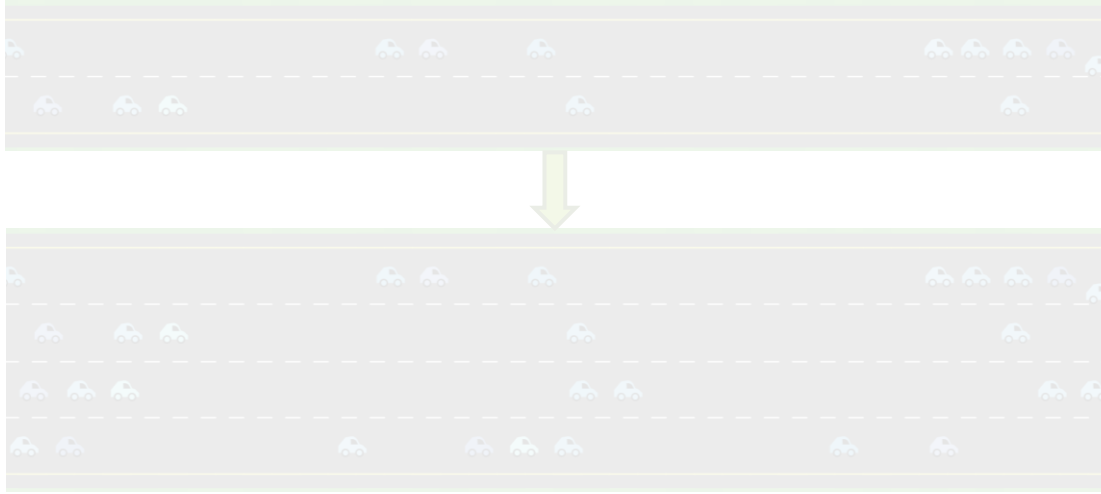
Equal chances of lane change ←

Constant traffic density ←

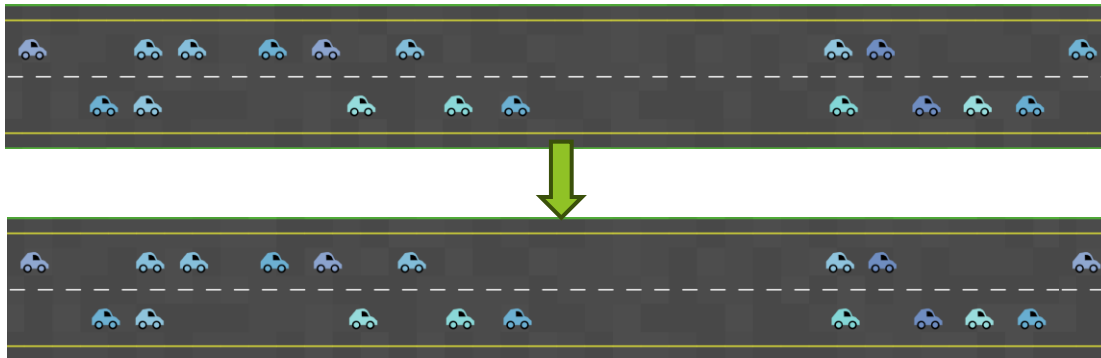
Scale both #cars and #lanes

# Intro - Scaling

► ~~Scale up #lanes~~



► Scale up #cars



Fair benchmarking ←

Equal chances of lane change ←

Constant traffic density ←

Scale #cars and lane\_length

Simplified for GPU optimization

# Intro - How the model works

► for each simulation step:

tryLaneChange

- // tryLaneChange
  - /\*start the clock\*/
  - determineTargetPosition()
  - tryLaneChange()
  - /\*time the clock\*/



driveForward

- // driveForward
  - /\*start the clock\*/
  - resolveCollisionsPerLane()
  - updateActualPosition()
  - /\*time the clock\*/



# Intro - How the model works

non-parallelizable

parallelize cars

parallelize lanes

for each simulation step:

- ▶ // tryLaneChange
  - ▶ /\*start the clock\*/
  - ▶ determineTargetPosition()
  - ▶ tryLaneChange()
  - ▶ /\*time the clock\*/
- ▶ // driveForward
  - ▶ /\*start the clock\*/
  - ▶ resolveCollisionsPerLane()
  - ▶ updateActualPosition()
  - ▶ /\*time the clock\*/



# Implementations

- ▶ CPU
  - ▶ serial codes, same logic
- ▶ GPU -> CUDA thrust + manual malloc (hybrid) -> let's call it "GPU thrust"
  - ▶ `thrust::for_each(..., ..., DetermineTargetPosition());`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `thrust::for_each(..., ..., UpdateActualPosition());`
- ▶ GPU -> CUDA manual malloc (purely manual) -> let's call it "GPU manual"
  - ▶ `determineTargetPositionCUDA<<<1, #threads>>>(...);`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `updateActualPositionCUDA<<<1, #threads>>>(...);`



# Implementations

- ▶ CPU
  - ▶ serial codes, same logic
- ▶ GPU -> CUDA thrust + manual malloc (hybrid) -> let's call it "GPU thrust"
  - ▶ `thrust::for_each(..., ..., DetermineTargetPosition());`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `thrust::for_each(..., ..., UpdateActualPosition());`
- ▶ GPU -> CUDA manual malloc (purely manual) -> let's call it "GPU manual"
  - ▶ `determineTargetPositionCUDA<<<1, #threads>>>(...);`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `updateActualPositionCUDA<<<1, #threads>>>(...);`

default #threads = 4

fixed #threads for lanes = 2

# Implementations

- ▶ CPU
  - ▶ serial codes, same logic
- ▶ GPU -> CUDA thrust + manual malloc (hybrid) -> let's call it "GPU thrust"
  - ▶ `thrust::for_each(..., ..., DetermineTargetPosition());`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `thrust::for_each(..., ..., UpdateActualPosition());`
- ▶ GPU -> CUDA manual malloc (purely manual) -> let's call it "GPU manual"
  - ▶ `determineTargetPositionCUDA<<<1, #threads>>>(...);`
  - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
  - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
  - ▶ `updateActualPositionCUDA<<<1, #threads>>>(...);`

CUDA Nsight Compute profiling, #cars=500

Function Name	De	Duration (1.10943e+09)	Ru (1.	Compute Throughput	Memory Throughput	# Registers	Grid Size	Block Size
determineTargetPositionCUDA	...	0.06	0.05	0.13	24	1...	4...	
tryLaneChangeCUDA	...	11.74	0.91	0.91	48	1...	1...	
resolveCollisionsPerLaneCUDA	...	0.18	0.08	0.08	40	1...	2...	
updateActualPositionCUDA	...	0.05	0.04	0.12	30	1...	4...	


default #threads = 4

fixed #threads for lanes = 2

# Implementations

## ► CPU

- serial codes, same logic



## ► GPU -> CUDA thrust + manual malloc (hybrid) -> let's call it "GPU thrust"

- `thrust::for_each(..., ..., DetermineTargetPosition());`

- `tryLaneChangeCUDA<<<1, 1>>>(...);`

- `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`

- `thrust::for_each(..., ..., UpdateActualPosition());`



## ► GPU -> CUDA manual malloc (purely manual) -> let's call it "GPU manual"

- `determineTargetPositionCUDA<<<1, #threads>>>(...);`

- `tryLaneChangeCUDA<<<1, 1>>>(...);`

- `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`

- `updateActualPositionCUDA<<<1, #threads>>>(...);`



# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

GPU thrust vs CPU:

- tryLaneChange significantly optimized!!
- driveForward runs terribly longer!

# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

GPU thrust vs CPU:

- tryLaneChange significantly optimized!!
- driveForward runs terribly longer!

GPU manual vs CPU:


- both tryLaneChange and driveForward are optimized!!

# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

- CPU tryLaneChange:
- Traffic size increases → Runtime **explodes!**




# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

	# lane changes	3,555	4,260	6,135	16,229
Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

- CPU tryLaneChange:
- Traffic size increases → Runtime **explodes!**
  - Successful lane changes adds #instructions




# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

- CPU driveForward:
- Traffic size increases → Runtime **increases linearly!**






# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

- GPU (thrust and manual) tryLaneChange:  
and GPU manual driveForward:
- Traffic size increases → Runtime **remains stable!**




# Benchmarking - CPU vs GPU-thrust vs GPU-manual

all runtime measured in us

Implementation	Task	500 Cars	1000 Cars	2000 Cars	4000 Cars
CPU	tryLaneChange	309,860	1,267,526	5,507,635	23,597,986
	driveForward	1,060	2,061	4,480	9,199
GPU thrust	tryLaneChange	2,083	2,263	2,182	2,333
	driveForward	1,388,718	4,347,869	24,827,547	99,747,402
GPU manual	tryLaneChange	1,849	2,066	1,587	2,137
	driveForward	614	596	697	695

- GPU thrust driveForward:
- Traffic size increases → Runtime **explodes!**



# What I've learnt

- ▶ CPU with increasing traffic size:
  - ▶ tryLaneChange runtime **explodes!** ← more successful lane changes, more instructions
  - ▶ driveForward **runs linearly slower** ← #instructions proportional to problem size

# What I've learnt

- ▶ CPU with increasing traffic size:
  - ▶ tryLaneChange runtime **explodes!** ← more successful lane changes, more instructions
  - ▶ driveForward **runs linearly slower** ← #instructions proportional to problem size
- ▶ from CPU to GPU with increasing traffic size:
  - ▶ tryLaneChange (both thrust and manual versions) **get optimized and remain stable** ← parallelization
  - ▶ driveForward (manual version) **get optimized and remain stable** ← parallelization
    - ▶ `determineTargetPositionCUDA<<<1, #threads>>>(...);`
    - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
    - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
    - ▶ `updateActualPositionCUDA<<<1, #threads>>>(...);`

# What I've learnt


- ▶ CPU with increasing traffic size:
  - ▶ tryLaneChange runtime **explodes!** ← more successful lane changes, more instructions
  - ▶ driveForward **runs linearly slower** ← #instructions proportional to problem size
- ▶ from CPU to GPU with increasing traffic size:
  - ▶ tryLaneChange (both thrust and manual versions) **get optimized and remain stable** ← parallelization
  - ▶ driveForward (manual version) **get optimized and remain stable** ← parallelization
    - ▶ `determineTargetPositionCUDA<<<1, #threads>>>(...);`
    - ▶ `tryLaneChangeCUDA<<<1, 1>>>(...);`
    - ▶ `resolveCollisionsPerLaneCUDA<<<1, 2>>>(...);`
    - ▶ `updateActualPositionCUDA<<<1, #threads>>>(...);`
  - ▶ driveForward (thrust version) **runs much slower!** ... too much abstraction in thrust
    - ▶ `thrust::for_each(carsDevice.begin(), carsDevice.end(), UpdateActualPosition());`

# Benchmarking - GPU manual malloc - different #threads

all runtime measured in us

Tasks (fixed #cars=2000)	#threads=4	#threads=8	#threads=16	#threads=32	#threads=64	#threads=128
determineTargetPosition<<<1, #threads>>>	1,715	1,967	1,866	1,222	1,137	1,753
updateActualPosition<<<1, #threads>>>	309	311	280	246	246	282

determineTargetPosition: best #threads = 64  
updateActualPosition: Not affected by #threads



# What I've learnt

- ▶ With fixed #cars=2000, as #threads increases →
  - ▶ determineTargetPosition **optimized!** (till #threads=64) ←
    - ▶ Arithmetic instruction: `car.TargetPosition = car.Position + car.TargetSpeed;`
  - ▶ No effects on updateActualPosition ←
    - ▶ Too simple instruction: `car.Position = car.TargetPosition;`

The background features abstract, overlapping green geometric shapes, primarily triangles and polygons, in various shades of green, creating a modern and dynamic visual effect.

# Thank you!

Tommy Poek

47192085